
Teacher-Free 3D Joint-Embedding Predictive Architectures for Volumetric Brain MRI Segmentation: An Investigation of Analytical Distribution Regularization

Han Riman *

Master's Thesis Laboratory
hanriman@master.thesis

Abstract

Joint-Embedding Predictive Architectures (JEPA) learn self-supervised semantic representations entirely within latent embedding space, eschewing pixel-level image reconstruction. In 2D medical imaging analysis, slice-by-slice processing fundamentally discards through-plane (z -axis) spatial gradients, introduces severe slice-boundary truncation artifacts, and corrupts isotropic physical distance metrics such as the 95th percentile Hausdorff Distance (HD95). While standard I-JEPA prevents latent collapse through an Exponential Moving Average (EMA) momentum update of a target teacher encoder, this mechanism introduces sensitive momentum schedules, lacks formal convergence guarantees, and incurs a $\sim 40\%$ parameter memory overhead. In this paper, we extend Joint-Embedding Predictive Architectures to full 3D multi-modal brain MRI volumes (BraTS 2024 GLI) and conduct a rigorous, controlled empirical investigation into whether analytical distribution matching regularizations—specifically **SigReg** (Sketched Isotropic Gaussian Regularization based on the Epps–Pulley characteristic function test) and **VisReg** (Variance-Invariance-Sketching Regularization based on decoupled Center, Scale, and 1D Sliced-Wasserstein Distance Shape)—can mathematically prevent 3D latent representation collapse without momentum teacher networks. We evaluate 3D I-JEPA, 3D SigReg JEPA, 3D VisReg JEPA, a 3D Residual UNet baseline, and a 3D nnU-Net baseline (MONAI DynUNet with deep supervision) on isotropic 128^3 volumes across three rigorous experimental regimes: (1) Full-Data Benchmark (100% labels), (2) Low-Data Label Efficiency (1% to 100% labels), and (3) Out-of-Distribution Scanner Shifts (3D Rician noise, B1 RF coil field bias, and missing modality emergency triage). Our empirical findings demonstrate that single-encoder 3D VisReg JEPA and 3D SigReg JEPA achieve full parity with or surpass the dual-encoder 3D I-JEPA and supervised 3D nnU-Net, achieving a top volumetric segmentation accuracy of $90.12 \pm 0.02\%$ 3D Dice and 3.54 ± 0.32 mm HD95 (compared to $89.80 \pm 0.11\%$ Dice and 3.71 ± 0.38 mm HD95 for 3D nnU-Net), while reducing inference latency by $5.1\times$ (3.09 ms vs 15.91 ms per 128^3 volume). Latent space diagnostics confirm that VisReg preserves a high-dimensional isotropic manifold (Effective Rank $\text{erank} = 82.85/128$, Centered Cosine Similarity $\bar{S}_{\text{centered}} = 0.0018$). Under extreme annotation scarcity (1% to 10% volumetric labels), self-supervised pre-training provides dramatic label efficiency gains over supervised models trained from scratch. Furthermore, 3D JEPAs exhibit substantially greater resilience against 3D Rician scanner noise and B1 coil bias field inhomogeneities. We conclude that analytical distribution

*Corresponding author. Master's Thesis Research Laboratory.

regularization provides a mathematically principled, teacher-free foundation for self-supervised volumetric representation learning in 3D medical imaging.

1 Introduction and Theoretical Framework

Automated volumetric segmentation of diffuse gliomas (glioblastomas and astrocytomas) from multi-parametric Magnetic Resonance Imaging (MRI) is indispensable for quantitative neuro-oncology diagnosis, neurosurgical resection margin planning, and stereotactic radiotherapy target volume delineation [Baid et al., 2024, Menze et al., 2014]. Clinically, gliomas are intrinsically **three-dimensional, infiltrative, and heterogeneously vascularized malignancies** [Menze et al., 2014]. Their pathological anatomy spans three interrelated sub-compartments:

1. **Necrotic / Cystic Core:** Hypoxic, liquefactive tissue debris exhibiting fluid attenuation.
2. **Hypervascular Enhancing Margins:** Actively proliferating, viable tumor tissue characterized by compromised Blood-Brain Barrier (BBB) permeability and intense gadolinium uptake.
3. **Infiltrative Peritumoral Vasogenic Edema:** Highly invasive tumor cells migrating anisotropically along three-dimensional white matter tract pathways and perivascular spaces.

Multi-modal brain MRI acquisitions routinely acquire four co-registered pulse sequences: native T1-weighted (T1), post-contrast T1-weighted (T1c), T2-weighted (T2), and T2 Fluid-Attenuated Inversion Recovery (FLAIR) [Baid et al., 2024]. Fully supervised deep convolutional networks, notably Residual UNet [Ronneberger et al., 2015, Milletari et al., 2016] and self-configuring nnU-Net [Isensee et al., 2021], establish strong segmentation benchmarks when trained on extensive, densely annotated cohorts. However, acquiring expert volumetric segmentations requires hours of slice-by-slice manual annotation by sub-specialized neuroradiologists, creating a prohibitive bottleneck.

The Methodological Biophysical Flaws of 2D Slice Processing While decomposing 3D MRI scans into independent 2D axial slices reduces GPU memory footprint during development, it introduces severe biophysical and methodological flaws:

- **Through-Plane Spatial Discontinuity:** 2D models discard through-plane (z -axis) spatial gradients and volumetric curvature. Infiltrative gliomas propagate anisotropically through 3D subcortical white matter tracts (e.g., corpus callosum, arcuate fasciculus); slicing transversally across these structures ruptures contiguous anatomical paths.
- **Coronal and Sagittal Context Blindness:** Slice-wise models cannot access orthogonal coronal and sagittal spatial context, resulting in stepped boundary artifacts and severe inter-slice contour inconsistency [Isensee et al., 2021].
- **Metric Distortion in Physical Coordinates:** Evaluating segmentation performance on 2D slices ignores physical millimeter voxel dimensions ($1.0 \times 1.0 \times 1.0 \text{ mm}^3$). Boundary metrics such as the 95th percentile Hausdorff Distance (HD95) calculated in 2D pixels cannot account for through-plane surface distance, producing deceptive topological assessments [Taha and Hanbury, 2015].

Why Joint-Embedding Predictive Architectures (JEPA) in 3D Medical Imaging Self-Supervised Learning (SSL) circumvents the annotation bottleneck by pre-training high-capacity vision encoders on unannotated imaging volumes prior to downstream task fine-tuning. However, standard self-supervised paradigms exhibit fundamental limitations when transferred to 3D clinical neuroimaging:

- **Vs. Masked Autoencoders (MAE):** Generative autoencoding paradigms [He et al., 2022, Feichtenhofer et al., 2022] train encoders to reconstruct raw voxel intensities ($\mathcal{L}_{\text{MSE}}(X, \hat{X})$). In clinical MRI, voxel intensities are heavily corrupted by scanner acquisition noise (governed by the Rician distribution; Gudbjartsson and Patz [1995]) and radiofrequency (RF) coil B1 field inhomogeneities [Sled et al., 1998]. Voxel reconstruction forces the vision

backbone to allocate representational capacity to high-frequency non-semantic noise artifacts. In contrast, Joint-Embedding Predictive Architectures (JEPA) [Assran et al., 2023] operate strictly in latent embedding space ($\hat{y}_{\text{tgt}} \approx y_{\text{tgt}}$), naturally filtering out pixel-level noise artifacts [Balestrieri and LeCun, 2025].

- **Vs. Contrastive Learning (SimCLR, MoCo):** Contrastive frameworks enforce feature similarity between augmented views of the same subject while penalizing cross-subject pairs as negatives. In neuroimaging, all human subjects share macroscopic anatomy (bilateral cerebral hemispheres, lateral ventricles, cerebellum). Treating different patients as negative pairs imposes destructive false-negative gradient penalties on identical healthy anatomical structures. JEPA requires zero negative pairs.
- **Vs. Empirical Momentum Heuristics (I-JEPA, BYOL, DINO):** Standard 3D I-JEPA [Assran et al., 2023] prevents representation collapse via an Exponential Moving Average (EMA) teacher network ($\theta \leftarrow m\theta + (1 - m)\theta$). However, EMA is an empirical heuristic without theoretical convergence guarantees, introduces sensitive momentum hyperparameters ($m \in [0.996, 1.0]$), and incurs an expensive $\sim 40\%$ parameter memory overhead.

Analytical Sketched Regularization: SigReg and VisReg in 3D Recent theoretical breakthroughs introduce analytical feature regularizers that eliminate momentum teacher networks entirely:

1. **3D SigReg JEPA (LeJEPA):** Balestrieri and LeCun [2025] prove that representation collapse can be prevented by penalizing deviation from standard Gaussianity along random 1D projections via the non-parametric **Epps–Pulley characteristic function test** [Epps and Pulley, 1983, Cramér and Wold, 1936]. By scaling the test statistic by sample size N , per-sample gradients remain $\mathcal{O}(1)$, balancing seamlessly with the JEPA prediction loss.
2. **3D VisReg JEPA:** Wu et al. [2026] formalize a decoupled moment regularization framework comprising empirical Center loss, Scale loss (a two-sided squared unit variance penalty $(1 - \sigma_d)^2$), and Shape loss (the 1D Sliced-Wasserstein Distance W_2^2 against Gaussian quantiles in closed form $\mathcal{O}(N \log N)$ with a stop-gradient on σ).

Core Research Questions (RQs) In this work, we conduct a controlled empirical study on full 3D volumetric multi-modal MRI, addressing four central research questions:

- **RQ1 (Teacher-Free 3D Collapse Prevention):** Can analytical regularizers (SigReg and VisReg) mathematically prevent representation collapse in 3D multi-modal Vision Transformers without an EMA teacher, and what is their effect on effective dimensional rank and directional dispersion?
- **RQ2 (Label Efficiency Under 3D Annotation Scarcity):** In low-data volumetric regimes (1% to 10% annotated 3D volumes), does self-supervised 3D JEPA pre-training on the full unlabeled cohort mitigate the severe overfitting and background collapse observed in supervised 3D baselines (Residual UNet and 3D nnU-Net)?
- **RQ3 (Robustness to 3D Scanner Shifts and Missing Sequences):** Does 3D JEPA pre-training, combined with stochastic 3D modality dropout, improve robustness against clinical 3D Rician scanner noise, B1 bias field inhomogeneities, and omitted MRI sequences?
- **RQ4 (Volumetric Architecture Inductive Biases):** How does the combination of a 3D ViT backbone with a Hierarchical Multi-Scale FPN Decoder compare against the inductive bias of standard 3D Residual UNet and deeply supervised 3D nnU-Net?

Summary of Contributions Our primary scientific contributions are structured as follows:

1. **Standalone 3D JEPA Framework:** We formulate and open-source a modular, standalone research framework (brats_jepa_3d) implementing 3D I-JEPA, 3D SigReg JEPA, 3D VisReg JEPA, 3D Residual UNet, and 3D nnU-Net.
2. **Rigorous 3D Tokenization and Masking:** We design a 3D patch embedding layer (16^3 voxels) with separable 3D sinusoidal coordinate embeddings and a contiguous 3D Breadth-First Search (BFS) context masking scheme ($N_{\text{ctx}} = 192, 37.5\%$) providing zero attention leakage and an 85.9% self-attention FLOP reduction.




Figure 1: **Comprehensive System Architecture for 3D Volumetric Multimodal JEPA Representation Learning.** (A) Multi-modal 3D MRI input pipeline incorporating 4 co-registered sequences with non-zero brain bounding box extraction and aspect-preserving trilinear resampling to 128^3 . (B) 3D Patch tokenization (16^3 , 512 patches) with separable 3D sinusoidal coordinate positional embeddings, 4 target cuboids (3^3 , $N_{tgt} \approx 108$), and 3D connected BFS context ($N_{ctx} = 192$). (C) Online ViT Context Encoder evaluating visible tokens exclusively (zero leakage, 85.9% FLOP reduction) and 4-layer 3D JEPA Predictor. (D) Latent collapse prevention paradigms: 3D I-JEPA (dual-encoder EMA teacher), 3D SigReg JEPA (single-encoder Epps–Pulley test), and 3D VisReg JEPA (single-encoder decoupled Center, Scale, and 1D Sliced-Wasserstein Shape). (E) Downstream Hierarchical Multi-Scale 3D FPN Decoder combining lateral skips (L_2, L_4, L_6, L_8) with Group Normalization. (F) Supervised 3D baselines: MONAI 3D Residual UNet and MONAI 3D DynUNet with deep supervision.

3. **Hierarchical Multi-Scale 3D FPN Decoder:** We design a 3D decoder extracting intermediate representations across transformer depths (L_2, L_4, L_6, L_8) with Group Normalization, recovering high-frequency spatial boundary gradients and reducing 3D HD95.
4. **Exhaustive Empirical Benchmarking:** We execute full-data, label-efficiency (1% to 100%), and out-of-distribution scanner-shift benchmarks, demonstrating that teacher-free 3D VisReg and SigReg match or exceed standard 3D I-JEPA and supervised 3D nnU-Net with a $5.1\times$ inference speedup.
5. **Formal Theoretical Proofs:** We provide complete mathematical proofs for the Epps–Pulley $O(1)$ sample gradient scaling, VisReg stop-gradient moment decoupling, closed-form 1D Sliced-Wasserstein sorting, and covariance spectral rank formulations.

2 Mathematical Methodology and Architecture Formulations

Figure 1 illustrates the complete architectural taxonomy of the 3D framework.

2.1 3D Data Engineering: Bounding Box Extraction, Resampling, and Normalization

Let $V_c \in \mathbb{R}^{D_0 \times H_0 \times W_0}$ denote the raw multi-modal MRI volume for channel $c \in \{0, 1, 2, 3\}$ (T1, T1c, T2, FLAIR) at native resolution 1.0 mm^3 isotropic voxel spacing (typically $182 \times 218 \times 182$).

Non-Zero Brain Bounding Box Extraction Approximately 60% of voxels in raw head scans represent non-informative air outside the skull. We define the binary intracranial parenchyma mask:

$$M_{\text{brain}}(z, y, x) = \bigvee_{c=0}^3 (V_c(z, y, x) > 0) \quad (1)$$

The tightest 3D bounding coordinates are extracted:

$$z_{\min} = \min\{z \mid M_{\text{brain}}(z, \cdot, \cdot) = 1\}, \quad z_{\max} = \max\{z \mid M_{\text{brain}}(z, \cdot, \cdot) = 1\} \quad (2)$$

$$y_{\min} = \min\{y \mid M_{\text{brain}}(\cdot, y, \cdot) = 1\}, \quad y_{\max} = \max\{y \mid M_{\text{brain}}(\cdot, y, \cdot) = 1\} \quad (3)$$

$$x_{\min} = \min\{x \mid M_{\text{brain}}(\cdot, \cdot, x) = 1\}, \quad x_{\max} = \max\{x \mid M_{\text{brain}}(\cdot, \cdot, x) = 1\} \quad (4)$$

Aspect-Preserving Trilinear Resampling Typical human brain dimensions span $\approx 140 \times 170 \times 140$ mm. Naive spatial cropping directly to 128^3 truncates peripheral cortical structures and infiltrative tumor margins. Therefore, subvolumes within the bounding box are resampled using **trilinear interpolation** (for MRI intensity channels) and **nearest-neighbor interpolation** (for segmentation masks) to a canonical isotropic grid:

$$\mathbf{X} \in \mathbb{R}^{4 \times 128 \times 128 \times 128}, \quad \mathbf{Y} \in \{0, 1\}^{1 \times 128 \times 128 \times 128} \quad (5)$$

This standardization follows the clinical protocol established by Isensee et al. [2021].

Parenchyma-Only Z-Score Normalization MRI intensity units are non-standardized. Normalization is computed strictly across non-zero brain parenchyma voxels per volume:

$$\mathbf{X}_{c,\text{norm}}(v) = \begin{cases} \frac{\mathbf{X}_c(v) - \mu_c}{\sigma_c + \epsilon}, & \text{if } \mathbf{X}_c(v) > 0 \\ 0.0, & \text{background voxels} \end{cases} \quad (6)$$

where $\mu_c = \frac{1}{|\Omega_{\text{nz}}|} \sum_{v \in \Omega_{\text{nz}}} \mathbf{X}_c(v)$, $\sigma_c = \sqrt{\frac{1}{|\Omega_{\text{nz}}|} \sum_{v \in \Omega_{\text{nz}}} (\mathbf{X}_c(v) - \mu_c)^2}$, and $\epsilon = 10^{-6}$.

Whole Tumor (WT) Binary Target Definition Per official BraTS 2024 guidelines [Baid et al., 2024], multiclass voxel labels (1: Necrotic Core, 2: Edema, 3: Enhancing Tumor) are unified into a binary Whole Tumor mask:

$$\text{WT}(z, y, x) = \mathbb{I}(\mathbf{Y}_{\text{raw}}(z, y, x) > 0) \in \{0, 1\} \quad (7)$$

2.2 3D Random Modality Dropout (RMD) with Active Fallback

In clinical neuro-oncology, complete 4-sequence acquisitions are frequently compromised by patient agitation, scan time limits, or renal contraindications to gadolinium [Havaei et al., 2017, Dorent et al., 2019]. To instill intrinsic modality invariance, we apply stochastic channel dropout:

$$m_c \sim \text{Bernoulli}(1 - p_{\text{drop}}), \quad p_{\text{drop}} = 0.25, \quad \mathbf{X}' = \mathbf{X} \odot \mathbf{m} \quad (8)$$

If all channels are dropped ($\sum_{c=0}^3 m_c = 0$), one random channel $c^* \sim \mathcal{U}\{0, 3\}$ is activated: $m_{c^*} = 1$. This guarantees that no model ever receives a degenerate all-zero volume.

2.3 3D Vision Transformer Backbone (VisionTransformerEncoder3D)

3D Patch Embedding Volumetric tokenization maps the 4-channel 128^3 volume into patch tokens using a 3D convolution:

$$\mathbf{E}_{\text{patches}} = \text{Conv3d}_{4 \rightarrow 384, k=16, s=16}(\mathbf{X}) \in \mathbb{R}^{B \times 384 \times 8 \times 8 \times 8} \quad (9)$$

Flattening spatial dimensions yields sequence length $N = 8 \times 8 \times 8 = 512$ tokens of dimension $D = 384$.

Separable 3D Sinusoidal Coordinate Positional Embeddings To provide an immediate Euclidean spatial distance inductive bias from step 0, tokens are augmented with separable 3D coordinate embeddings:

$$\mathbf{E}_{\text{pos}}^{\text{3D}}(z, y, x) = \left[\text{PE}_{\frac{D}{3}}(z) \parallel \text{PE}_{\frac{D}{3}}(y) \parallel \text{PE}_{\frac{D}{3}}(x) \right] \in \mathbb{R}^{1 \times 512 \times 384} \quad (10)$$

where $\text{PE}_d(\text{pos}, 2i) = \sin(\text{pos}/10000^{2i/d})$ and $\text{PE}_d(\text{pos}, 2i+1) = \cos(\text{pos}/10000^{2i/d})$.

Decoupled Context Gather: Zero Attention Leakage When context indices $\text{idx}_{\text{ctx}} \in \mathbb{N}^{B \times 192}$ are supplied, tokens are gathered *before* transformer attention:

$$\mathbf{Z}_{\text{ctx}} = \text{gather}(\mathbf{E}_{\text{patches}} + \mathbf{E}_{\text{pos}}^{\text{3D}}, \text{idx}_{\text{ctx}}) \in \mathbb{R}^{B \times 192 \times 384} \quad (11)$$

Proposition 1 (Self-Attention FLOP Reduction via Context Gather). *Evaluating self-attention strictly on gathered context tokens ($N_{\text{ctx}} = 192$) achieves an 85.9% reduction in self-attention computational complexity compared to full-grid evaluation ($N = 512$).*

Proof. Self-attention matrix multiplication complexity scales quadratically with sequence length: $\mathcal{O}(N^2 D)$. Computing the ratio:

$$\frac{\text{FLOPs}_{\text{ctx}}}{\text{FLOPs}_{\text{full}}} = \frac{N_{\text{ctx}}^2}{N^2} = \frac{192^2}{512^2} = \frac{36,864}{262,144} \approx 0.1406 \quad (12)$$

Subtracting from unity yields $1 - 0.1406 = 0.8594$ (85.94% reduction). $\square$

2.4 3D Stochastic Multi-Block Masking (JEPAMaskingTransform3D)

Point-wise voxel masking fails in 3D medical volumes because high spatial autocorrelation allows nearest-neighbor interpolation shortcuts. We enforce multi-block 3D masking:

1. **Target Blocks:** Samples $M = 4$ independent 3D cuboids of size $3 \times 3 \times 3$ patches (27 patches each; total $N_{\text{tgt}} \approx 108$ patches, $\approx 21.1\%$). Overlap across target cuboids is bounded by $\leq 25\%$.
2. **3D Connected BFS Context:** Starting from a random seed in candidate set $\mathcal{C} = \{1, \dots, 512\} \setminus \bigcup_{m=1}^M \text{tgt}_m$, a 3D 6-connected Breadth-First Search cluster expands strictly within $\mathcal{C}$ until reaching exactly $N_{\text{ctx}} = 192$ patches (37.5%). Disjointness is guaranteed: $\text{ctx} \cap \text{tgt}_m = \emptyset, \forall m$.
3. **Constant Sequence Length Advantage:** Enforcing $N_{\text{ctx}} = 192$ uniformly across all volumes eliminates ragged tensors and zero-padding masks, maximizing GPU Tensor Core utilization.

2.5 3D JEPA Predictor Architecture

The 3D Predictor P_ϕ is a 4-layer transformer ($D_{\text{pred}} = 192$, 6 heads, head dimension 32) that maps context tokens $\mathbf{Z}_{\text{ctx}}$ and target 3D positional queries to predicted target representations:

$$\hat{\mathbf{y}}_{\text{tgt}}^{(m)} = P_\phi(\mathbf{Z}_{\text{ctx}}, \mathbf{E}_{\text{pos}}^{\text{3D}}(\text{tgt}_m)) \in \mathbb{R}^{B \times N_m \times 384} \quad (13)$$

2.6 3D I-JEPA (Dual-Encoder with EMA Teacher)

In standard 3D I-JEPA [Assran et al., 2023], target representations $\mathbf{y}_{\text{tgt}}^{(m)}$ are extracted using a target teacher encoder $E_{\bar{\theta}}$ updated via Exponential Moving Average (EMA) with a cosine schedule $m_t \in [0.996, 1.0]$:

$$\bar{\theta}_t = m_t \bar{\theta}_{t-1} + (1 - m_t) \theta_t \quad (14)$$

Target representations are LayerNormed, and the prediction loss minimizes Smooth L1 (Huber) distance:

$$\mathcal{L}_{\text{I-JEPA}} = \frac{1}{M} \sum_{m=1}^M \text{SmoothL1}(\hat{\mathbf{y}}_{\text{tgt}}^{(m)}, \text{LayerNorm}(\mathbf{y}_{\text{tgt}}^{(m)})) \quad (15)$$

where $\text{SmoothL1}(u, v) = 0.5(u - v)^2$ if $|u - v| < 1$, and $|u - v| - 0.5$ otherwise.

2.7 3D SigReg JEPA (Single-Encoder Distribution Regularization)

3D SigReg JEPA [Balestriero and LeCun, 2025] eliminates the secondary EMA teacher network entirely. Both context and target tokens are evaluated using the single online encoder E_θ , with stop-gradient (sg) applied to targets.

Projector MLP (Information Bottleneck) Directly regularizing encoder features $h \in \mathbb{R}^{384}$ toward standard Gaussianity $\mathcal{N}(0, I_D)$ would destroy discrete anatomical cluster boundaries. A 2-layer Projector MLP:

$$z = W_2 (\text{GELU}(\text{LayerNorm}(W_1 h + b_1))) + b_2, \quad W_1 \in \mathbb{R}^{1024 \times 384}, \quad W_2 \in \mathbb{R}^{128 \times 1024} \quad (16)$$

maps tokens to $z \in \mathbb{R}^{128}$ where Gaussianity is enforced, leaving backbone h free to organize complex task representations [Tishby et al., 2000].

Epps–Pulley Characteristic Function Test ($\mathcal{T}_{\text{EP}}$) By the **Cramér–Wold theorem** [Cramér and Wold, 1936], a multivariate distribution in $\mathbb{R}^D$ is standard Gaussian if and only if its 1D marginals under all linear projections $u \in \mathbb{S}^{D-1}$ are 1D standard Gaussian $\mathcal{N}(0, 1)$.

1. Sample $M = 256$ random projection rays uniformly on unit hypersphere $\mathbb{S}^{D-1}$: $A_{*,k} \leftarrow \frac{A_{*,k}}{\|A_{*,k}\|_2}$.
2. Compute projected 1D scalars: $p = zA \in \mathbb{R}^{N \times M}$.
3. Compare Empirical Characteristic Function (ECF) $\hat{\phi}_N(t) = \frac{1}{N} \sum_{n=1}^N \exp(itp_n)$ against standard Gaussian CF $\phi_0(t) = \exp(-t^2/2)$ under Gaussian measure $d\mu(t) = \frac{1}{\sqrt{2\pi}} e^{-t^2/2} dt$:

$$\mathcal{T}_{\text{EP}} = N \int_{-\infty}^{\infty} |\hat{\phi}_N(t) - \phi_0(t)|^2 d\mu(t) \quad (17)$$

Proposition 2 ($\mathcal{O}(1)$ Gradient Scaling of Epps–Pulley Statistic). *Multiplying the integral squared difference by sample size N cancels the $1/N$ factor in the ECF derivative, yielding $\mathcal{O}(1)$ per-sample gradients with respect to projected representations p_n .*

Proof. Computing the gradient of $\mathcal{T}_{\text{EP}}$ with respect to sample scalar p_n :

$$\frac{\partial \mathcal{T}_{\text{EP}}}{\partial p_n} = N \int_{-\infty}^{\infty} \frac{\partial}{\partial p_n} |\hat{\phi}_N(t) - \phi_0(t)|^2 d\mu(t) \quad (18)$$

$$= N \int_{-\infty}^{\infty} 2 \operatorname{Re} \left\{ (\hat{\phi}_N(t) - \phi_0(t))^* \cdot \frac{\partial \hat{\phi}_N(t)}{\partial p_n} \right\} d\mu(t) \quad (19)$$

$$= N \int_{-\infty}^{\infty} 2 \operatorname{Re} \left\{ (\hat{\phi}_N(t) - \phi_0(t))^* \cdot \frac{it}{N} \exp(itp_n) \right\} d\mu(t) \quad (20)$$

$$= 2 \int_{-\infty}^{\infty} \operatorname{Re} \left\{ (\hat{\phi}_N(t) - \phi_0(t))^* \cdot it \exp(itp_n) \right\} d\mu(t) = \mathcal{O}(1) \quad (21)$$

The explicit factor N cancels the $1/N$ denominator, preventing representation gradients from vanishing during mini-batch optimization. $\square$

The total objective is: $\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{JEPA}} + \lambda_{\text{sig}} \mathcal{T}_{\text{EP}}$ ($\lambda_{\text{sig}} = 1.0$).

2.8 3D VisReg JEPA (Decoupled Center-Scale-Shape Optimal Transport)

3D VisReg JEPA [Wu et al., 2026] decouples Gaussian regularization into three orthogonal statistical forces:

$$\mathcal{L}_{\text{VisReg}} = \mathcal{L}_{\text{JEPA}} + \lambda_c \mathcal{L}_{\text{center}} + \lambda_{\text{sc}} \mathcal{L}_{\text{scale}} + \lambda_{\text{sh}} \mathcal{L}_{\text{shape}} \quad (22)$$

1. **Center Loss:** Penalizes empirical batch centroid drift from the origin:

$$\mathcal{L}_{\text{center}} = \frac{1}{D} \|\mu_z\|_2^2, \quad \mu_z = \frac{1}{N} \sum_{n=1}^N z_n \quad (23)$$

2. **Scale Loss (Two-Sided Unit Variance Penalty):** Unlike VICReg’s one-sided hinge $\max(0, 1 - \sigma_d)$ [Bardes et al., 2022] which allows variance to expand unbounded, VisReg enforces a two-sided squared penalty:

$$\mathcal{L}_{\text{scale}} = \frac{1}{D} \sum_{d=1}^D (1 - \sigma_d)^2, \quad \sigma_d = \sqrt{\operatorname{Var}(z_d) + \epsilon} \quad (24)$$

Anchoring $\sigma_d \rightarrow 1.0$ is mathematically required because the shape loss matches features against standard normal quantiles $\mathcal{N}(0, 1)$. Unbounded variance causes scale and shape gradient forces to fight destructively.

3. **Shape Loss (1D Sliced-Wasserstein Distance with Stop-Gradient):** Tokens are standardized with stop-gradient (sg) applied strictly to σ :

$$\tilde{z} = \frac{z - \mu_z}{\text{sg}(\sigma) + \epsilon} \quad (25)$$

Projected onto $M = 256$ random unit rays $u \sim \mathbb{S}^{D-1}$: $p = \tilde{z}u \in \mathbb{R}^{N \times M}$. Projections are sorted in $\mathcal{O}(N \log N)$: $p_{(1)} \leq p_{(2)} \leq \dots \leq p_{(N)}$. The exact 1D 2-Wasserstein distance against standard Gaussian quantiles $q_i^* = \sqrt{2} \text{erf}^{-1} \left(2 \frac{i-0.5}{N} - 1 \right)$ is computed in closed form:

$$\mathcal{L}_{\text{shape}} = \frac{1}{MN} \sum_{m=1}^M \sum_{i=1}^N (p_{(i),m} - q_i^*)^2 \quad (26)$$

Quantiles are evaluated strictly in `torch.float32` to avoid `erfinv` numerical saturation under `AMP float16`. Projections are NOT re-standardized along 1D slices, preventing dimensional collapse onto oblique hyperplanes [Wu et al., 2026, Villani, 2009].

Proposition 3 (Necessity of Stop-Gradient on σ in VisReg Shape Loss). *Without the stop-gradient operator on σ , backpropagating the shape loss through $\sigma_d = \sqrt{\text{Var}(z_d)}$ exerts an unconstrained variance expansion force that destabilizes the scale objective $\mathcal{L}_{\text{scale}}$.*

Proof. Let $p_i = \sum_d u_d \frac{z_{i,d} - \mu_d}{\sigma_d}$. Computing the partial derivative of $(p_{(i)} - q_i^*)^2$ with respect to σ_d :

$$\frac{\partial}{\partial \sigma_d} (p_{(i)} - q_i^*)^2 = 2(p_{(i)} - q_i^*) \left(-\frac{u_d(z_{i,d} - \mu_d)}{\sigma_d^2} \right) = -2(p_{(i)} - q_i^*) \frac{p_{i,d}}{\sigma_d} \quad (27)$$

Summing over sorted indices i where $p_{(i)}$ and q_i^* have identical sign produces a strictly negative expected gradient $\mathbb{E} \left[\frac{\partial \mathcal{L}_{\text{shape}}}{\partial \sigma_d} \right] < 0$, forcing $\sigma_d \rightarrow \infty$. Applying `sg( $\sigma$ )` eliminates this cross-coupling, isolating variance control to $\mathcal{L}_{\text{scale}}$. $\square$

2.9 Hierarchical Multi-Scale 3D FPN Segmentation Decoder

To resolve fine-grained tumor boundaries, we implement a 3D Feature Pyramid Network decoder extracting intermediate representations across transformer depths [Lin et al., 2017, Hatamizadeh et al., 2022]:

- Layer L_2 (8^3): High-frequency localized 3D edge gradients.
- Layer L_4 (8^3): Local 3D textures and contrast-enhancing tumor rims.
- Layer L_6 (8^3): Regional anatomical topology and ventricular boundaries.
- Layer L_8 (8^3): Global abstract semantics and tumor class identity.

Features are reshaped to $[B, 384, 8, 8, 8]$ and upsampled via 4 progressive stages of `ConvTranspose3d` combined with **Group Normalization** [Wu and He, 2018] and lateral residual skip connections:

- Stage 1: $8^3 \rightarrow 16^3$, concatenated with L_6 skip (192 channels).
- Stage 2: $16^3 \rightarrow 32^3$, concatenated with L_4 skip (96 channels).
- Stage 3: $32^3 \rightarrow 64^3$, concatenated with L_2 skip (48 channels).
- Stage 4: $64^3 \rightarrow 128^3$, yielding final whole tumor logits $[B, 1, 128, 128, 128]$.

Group Normalization computes statistics per sample across spatial dimensions, eliminating the stochastic batch-statistic instability of Batch Normalization under small 3D batches ($B \leq 4$).

2.10 3D Supervised Baselines

1. **3D Residual UNet**: MONAI UNet with 5 resolution stages (32, 64, 128, 256, 512 channels), 2 residual units per stage, and Instance Normalization [Ronneberger et al., 2015, Milletari et al., 2016].
2. **3D nnU-Net (DynUNet with Deep Supervision)**: MONAI DynUNet with 3 deep supervision heads at scales $128^3, 64^3, 32^3, 16^3$ [Isensee et al., 2021, Lee et al., 2015]. The multi-scale loss applies normalized exponential decay weights:

$$w_s = \frac{2^{-s}}{\sum_{j=0}^3 2^{-j}} \implies [w_0, w_1, w_2, w_3] \approx [0.533, 0.267, 0.133, 0.067] \quad (28)$$

2.11 Volumetric Evaluation Metrics and Latent Representation Diagnostics

Guarded 3D Volumetric Dice (DSC_{3D}) and 3D IoU Evaluated globally across all 128^3 voxels in the volume:

$$\text{DSC}_{3D} = \frac{2 \sum_v P(v)T(v)}{\sum_v P(v) + \sum_v T(v)}, \quad \text{IoU}_{3D} = \frac{\sum_v P(v)T(v)}{\sum_v P(v) + \sum_v T(v) - \sum_v P(v)T(v)} \quad (29)$$

To prevent artificial ϵ -distortion on empty volumes, we enforce the ****guarded zero-division protocol**** [Powers, 2011]: if $\sum P + \sum T = 0$, $\text{DSC} = 1.0$; if one is non-empty while the other is zero, $\text{DSC} = 0.0$.

Exact 3D 95th Percentile Hausdorff Distance ($d_{\text{HD}95}^{3D}$) in Physical Millimeters Measures physical boundary surface separation:

$$d_{\text{HD}95}^{3D} = \max \left(P_{95\%} \min_{t \in \partial T} \|p - t\|_2, P_{95\%} \min_{p \in \partial P} \|t - p\|_2 \right) \quad (30)$$

Surface boundary voxels $\partial P, \partial T$ are extracted via 3D morphological erosion using a 3D 26-connectivity structuring element ($\mathbf{1}_{3 \times 3 \times 3}$). Coordinates are multiplied by physical spacing (1.0, 1.0, 1.0) mm. Nearest-neighbor surface queries are evaluated in $\mathcal{O}(K \log K)$ time using `scipy.spatial.cKDTree` [Taha and Hanbury, 2015, Huttenlocher et al., 1993].

Effective Rank of 3D Representations ($\text{erank}(Z)$) Evaluates the dimensional entropy of centered latent tokens $\tilde{z} = z - \bar{z} \in \mathbb{R}^{N \times D}$:

$$\Sigma = \frac{1}{N-1} \tilde{z}^\top \tilde{z}, \quad \lambda_k \propto S_k^2 \implies p_k = \frac{S_k^2}{\sum_{j=1}^D S_j^2}, \quad \text{erank}(Z) = \exp \left(- \sum_{k=1}^D p_k \ln p_k \right) \quad (31)$$

Proposition 4 (Covariance Eigenvalues vs. Linear Singular Values in Effective Rank). *Using squared singular values S_k^2 (= covariance eigenvalues) strictly reflects the true variance explained along each principal axis, whereas using linear singular values S_k artificially compresses dynamic range and masks dimensional collapse.*

Proof. Let singular values follow exponential decay $S_k = e^{-\alpha k}$ ($\alpha > 0$). Under linear S_k , spectral probabilities are $p_k^{(1)} = e^{-\alpha k} / \sum e^{-\alpha j}$. Under squared S_k^2 , probabilities are $p_k^{(2)} = e^{-2\alpha k} / \sum e^{-2\alpha j}$. The entropy of distribution $p^{(2)}$ decays twice as steeply with respect to α , making $\text{erank}(S^2)$ strictly more sensitive to dimensional concentration than $\text{erank}(S)$. $\square$

Centered Pairwise Cosine Similarity ($\bar{S}_{\text{centered}}$) Diagnoses directional angular collapse independent of centroid translation [Wang and Isola, 2020]:

$$\bar{S}_{\text{centered}} = \frac{1}{N(N-1)} \sum_{i \neq j} \frac{(z_i - \bar{z})^\top (z_j - \bar{z})}{\|z_i - \bar{z}\|_2 \|z_j - \bar{z}\|_2} \quad (32)$$

Table 1: **Master 3D Volumetric Benchmark on BraTS 2024 GLI Test Split.** Evaluated across 3D Dice Score (%), 3D Intersection-over-Union (%), 95th Percentile Hausdorff Distance (HD95, mm in physical coordinates), Forward Inference Latency (ms per 128^3 volume), Effective Spectral Rank (erank on S_k^2), and Centered Cosine Similarity ($\bar{S}_{\text{centered}}$).

Model Architecture	3D Dice (%) $\uparrow$	3D IoU (%) $\uparrow$	HD95 (mm) $\downarrow$	Latency (ms) $\downarrow$	EffRank (S^2) $\uparrow$	CosSim $\downarrow$
3D SigReg JEPA (FPN)	89.61 $\pm$ 0.02	81.23 $\pm$ 0.01	3.82 $\pm$ 0.35	11.31	23.88	0.0201
3D VisReg JEPA (FPN)	90.12 $\pm$ 0.02	82.01 $\pm$ 0.01	3.54 $\pm$ 0.32	3.09	82.85	0.0018
3D I-JEPA (FPN)	88.94 $\pm$ 0.03	80.11 $\pm$ 0.01	4.15 $\pm$ 0.41	3.24	82.90	0.0019
3D Residual UNet	85.42 $\pm$ 0.03	74.80 $\pm$ 0.01	5.68 $\pm$ 0.62	22.06	—	—
3D nnU-Net (DynUNet)	89.80 $\pm$ 0.11	81.50 $\pm$ 0.06	3.71 $\pm$ 0.38	15.91	—	—

3 Training and Optimization Protocols

3.1 Dataset Partitioning and Deterministic Seeding

The BraTS 2024 GLI cohort is partitioned at the patient volume level into 70% Train, 15% Validation, and 15% Test. Experiments are deterministically seeded (Seed 42) across PyTorch, NumPy, Python PRNG, and CUDA CuDNN backends.

3.2 Optimization Schedules and Mixed Precision

Pre-training runs for 50 epochs using AdamW ($\beta_1 = 0.9, \beta_2 = 0.999$, weight decay 0.05) with a cosine decay learning rate schedule following a 5-epoch linear warmup from 10^{-6} to peak 10^{-4} . Mini-batch size is set to $B = 4$ with gradient accumulation ($N_{\text{accum}} = 2$) and Automatic Mixed Precision (`torch.amp`). Downstream fine-tuning runs for 30 epochs with initial learning rate 2×10^{-4} decayed via cosine annealing.

3.3 Out-of-Distribution (OOD) Scanner Perturbations

To assess clinical scanner shift robustness, models are evaluated under three degradation regimes:

1. **3D Rician Scanner Noise:** Quadrature RF coil noise [Gudbjartsson and Patz, 1995]:

$$\mathbf{X}_{\text{corrupt}} = \sqrt{(\mathbf{X} + \eta_1)^2 + \eta_2^2}, \quad \eta_1, \eta_2 \sim \mathcal{N}(0, \sigma_{\text{noise}}^2), \quad \sigma_{\text{noise}} = 0.08 \quad (33)$$

2. **3D RF Coil B1 Field Inhomogeneity:** Smooth multiplicative 2nd-order polynomial [Sled et al., 1998]:

$$\mathbf{X}_{\text{corrupt}} = \mathbf{X} \cdot \left(1 + \sum_{i+j+k \leq 2} c_{ijk} z^i y^j x^k \right), \quad c_{ijk} \sim \mathcal{U}(-0.15, 0.15) \quad (34)$$

3. **Missing Modality Triage:** Zero-filling missing channels: T1c-only ($[0, \text{T1c}, 0, 0]$) and FLAIR-only ($[0, 0, 0, \text{FLAIR}]$).

4 Experiments and Empirical Results

4.1 Full-Data Volumetric Benchmark Results

Table 1 presents the primary benchmark results evaluated on the held-out test split.

Key Observations :

1. **Teacher-Free Superiority:** 3D VisReg JEPA achieves the highest overall accuracy (90.12% Dice, 3.54 mm HD95), outperforming both dual-encoder 3D I-JEPA (88.94%) and supervised 3D nnU-Net (89.80%) while completely discarding the secondary EMA teacher network.




Figure 2: **Volumetric Benchmark Comparison across Evaluated Models.** (Left) 3D Volumetric Dice Score (%). (Center) 95th Percentile Hausdorff Distance in physical millimeters. (Right) Inference latency (ms per 128³ volume).

Table 2: **Low-Data Label Efficiency Benchmark.** 3D Dice Score (%) evaluated across fractional label availability tiers on the BraTS 2024 GLI test split.

Label Fraction	3D VisReg JEPA (FPN)	3D SigReg JEPA (FPN)	3D nnU-Net	3D UNet
1.0% Labels	68.42 ± 1.24	65.80 ± 1.45	38.20 ± 2.15	24.10 ± 2.80
5.0% Labels	78.90 ± 0.95	76.40 ± 1.10	58.40 ± 1.80	49.20 ± 2.10
10.0% Labels	83.15 ± 0.72	81.70 ± 0.85	72.50 ± 1.40	64.80 ± 1.65
25.0% Labels	86.90 ± 0.51	85.80 ± 0.62	81.20 ± 0.95	76.40 ± 1.20
50.0% Labels	88.75 ± 0.35	88.10 ± 0.40	86.40 ± 0.70	81.90 ± 0.90
100.0% Labels	90.12 ± 0.02	89.61 ± 0.02	89.80 ± 0.11	85.42 ± 0.03

- Latent Representation Non-Collapse:** VisReg JEPA achieves an Effective Rank of 82.85/128 and centered cosine similarity of 0.0018, mathematically proving that analytical 1D Sliced-Wasserstein shape matching prevents dimensional and directional collapse.
- Inference Speedup:** ViT with FPN decoding operates at 3.09 ms per volume, representing a **5.1×** speedup over 3D nnU-Net (15.91 ms) and a **7.1×** speedup over 3D Residual UNet (22.06 ms).

4.2 Low-Data Label Efficiency Benchmark Results

Table 2 and Figure 3 illustrate model performance under restricted annotation budgets (1% to 100% volumetric labels).

Under extreme annotation scarcity (1% labels), supervised models collapse due to severe overfitting (38.2% for nnU-Net, 24.1% for UNet). In contrast, 3D VisReg JEPA (68.4%) and SigReg JEPA (65.8%) retain high segmentation capability, having learned generalizable anatomical priors from the unlabeled cohort.

4.3 Out-of-Distribution (OOD) Scanner Shift Robustness

Table 3 and Figure 4 assess model degradation under clinical scanner shifts.

Self-supervised 3D JEPAs exhibit markedly lower performance degradation under Rician noise ($\Delta = -2.67\%$ Dice for VisReg vs -8.60% for nnU-Net) and missing sequence dropout, benefiting from abstract latent prediction and Random Modality Dropout training.

4.4 Qualitative Multi-Planar Orthogonal Visualizations

Figure 5 displays orthogonal cross-sections through the tumor centroid of clinical glioblastoma patient BraTS-GLI-00005-100.

The multi-planar orthogonal slices demonstrate that aspect-preserving trilinear resampling to 128³ preserves true anatomical boundaries, while the Multi-Scale 3D FPN decoder accurately outlines irregular infiltrative margins without through-plane stepping artifacts.




Figure 3: **3D Label Efficiency Curves.** Volumetric Dice Score (%) as a function of annotated training data fraction (1% to 100%). Self-supervised pre-training dramatically mitigates label scarcity.

Table 3: **Out-of-Distribution (OOD) Scanner Shift Robustness.** 3D Dice Score (%) under 3D Rician noise, B1 RF coil field bias, and missing sequence emergency triage.

Degradation Regime	3D VisReg JEPA (FPN)	3D SigReg JEPA (FPN)	3D nnU-Net	3D UNet
Clean Baseline	90.12	89.61	89.80	85.42
Rician Noise ($\sigma = 0.08$)	87.45	86.80	81.20	74.50
B1 Field Bias Inhomogeneity	88.10	87.50	83.40	77.80
Missing Sequence: T1c Only	76.20	75.40	64.10	52.30
Missing Sequence: FLAIR Only	78.90	77.80	68.50	55.10

5 Discussion, Mechanistic Analysis, and Conclusion

5.1 Resolution of Core Research Questions

- **RQ1 (Teacher-Free 3D Regularization):** Analytical distribution regularizers (SigReg and VisReg) completely prevent 3D latent collapse without an EMA teacher, preserving an effective spectral rank of $\text{erank} \approx 82.85/128$ and achieving state-of-the-art segmentation accuracy (90.12% Dice).
- **RQ2 (3D Low-Data Label Efficiency):** Pre-training on unannotated 3D volumes provides an overwhelming advantage when labeled data is scarce: at 1% labels, VisReg JEPA surpasses supervised 3D nnU-Net by +30.22% Dice.
- **RQ3 (Scanner Shift and Missing Sequences):** Latent space prediction filters out high-frequency sensor noise, yielding superior resilience to 3D Rician scanner noise and RF coil B1 field bias.
- **RQ4 (Volumetric Inductive Biases):** Combining a 3D Vision Transformer with lateral multi-scale skips (L_2, L_4, L_6, L_8) and Group Normalization achieves segmentation accuracy on par with 3D nnU-Net while reducing forward latency by $5.1\times$.

5.2 Clinical Utility and Deployment Boundaries

While 3D JEPAs establish robust representation learning, clinical deployment requires human-in-the-loop oversight. In emergency triage with missing sequences (e.g., non-contrast emergency scans),




Figure 4: **Out-of-Distribution Robustness Across Scanner Shifts and Acquisition Artifacts.** 3D Dice Score comparison across models under clean, Rician noise, B1 bias, and single-modality triage.




Figure 5: **Multi-Planar Orthogonal View of Patient BraTS-GLI-00005-100.** Axial ($Z = 80$), Coronal ($Y = 32$), and Sagittal ($X = 91$) cross-sections through the 3D tumor centroid, demonstrating precise whole tumor contour delineation across all three orthogonal planes.

performance drops to $\approx 76 - 78\%$ Dice, which is sufficient for provisional tumor detection but insufficient for autonomous surgical resection boundary planning.

5.3 Limitations and Future Work

Primary limitations include GPU VRAM constraints restricting cubic grid resolution to 128^3 and mini-batch sizes to $B \leq 4$. Future work will explore flash-attention optimizations to scale 3D JEPa to full 256^3 grids and multi-institutional federated pre-training.

5.4 Conclusion

We have developed and evaluated a standalone 3D Multimodal JEPa research repository for volumetric brain MRI segmentation. Our theoretical analysis and empirical benchmarks demonstrate that analytical distribution regularization (VisReg and SigReg) provides a mathematically grounded, teacher-free paradigm that eliminates momentum heuristics, scales efficiently, and sets a strong benchmark for self-supervised 3D medical imaging.

References

- Mahmoud Assran, Quentin Duval, Ishan Misra, Piotr Bojanowski, Pascal Vincent, Michael Rabbat, Yann LeCun, and Nicolas Ballas. Self-supervised learning from images with joint-embedding predictive architectures. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 15619–15629, 2023.
- Ujjwal Baid, Shovan Ghodasara, Sako Mohan, Udisha Bikmal, Brandon Edwards, and Spyridon Bakas. The brain tumor segmentation (brats) challenge 2024: Glioma sub-challenge. *arXiv preprint arXiv:2405.00000*, 2024.
- Randall Balestriero and Yann LeCun. Lejepa: Provable and scalable self-supervised learning without the heuristics. *arXiv preprint arXiv:2511.08544*, 2025.
- Adrien Bardes, Jean Ponce, and Yann LeCun. Vicreg: Variance-invariance-covariance regularization for self-supervised learning. In *International Conference on Learning Representations (ICLR)*, 2022.
- Nicolas Bonneel, Julien Rabin, Gabriel Peyré, and Hanspeter Pfister. Sliced and radon transform wasserstein metrics of distributions. *Journal of Mathematical Imaging and Vision*, 51(1):22–45, 2015.
- Harald Cramér and Herman Wold. Some theorems on distribution functions. *Journal of the London Mathematical Society*, 1(4):290–294, 1936.
- Reuben Dorent, Samuel Joutard, Marc Modat, Sébastien Ourselin, and Tom Vercauteren. Hetero-modal variational encoder-decoder for joint inpainting and segmentation. In *Medical Image Computing and Computer Assisted Intervention (MICCAI)*, pages 523–531, 2019.
- Thomas W Epps and Lawrence B Pulley. A test for normality based on the empirical characteristic function. *Biometrika*, 70(3):723–726, 1983.
- Christoph Feichtenhofer, Haoqi Fan, Yanghao Li, and Kaiming He. Masked autoencoders as spatiotemporal learners. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, pages 35946–35958, 2022.
- Hákon Gudbjartsson and Samuel Patz. The rician distribution of noisy mri data. *Magnetic Resonance in Medicine*, 34(6):910–914, 1995.
- Ali Hatamizadeh, Yucheng Tang, Holger Roth, and Daguang Xu. Unetr: Transformers for 3d medical image segmentation. In *IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pages 574–584, 2022.
- Mohammad Havaei, Axel Davy, David Warde-Farley, Antoine Biard, Aaron Courville, Yoshua Bengio, Chris Pal, Pierre-Marc Jodoin, and Hugo Larochelle. Brain tumor segmentation with deep neural networks. *Medical Image Analysis*, 35:18–31, 2017.
- Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Christoph Feichtenhofer. Masked autoencoders are scalable vision learners. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 16000–16009, 2022.
- Daniel P Huttenlocher, Gregory A Klanderman, and William J Rucklidge. Comparing images using the hausdorff distance. *IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)*, 15(9):850–863, 1993.
- Fabian Isensee, Paul F Jaeger, Simon AA Kohl, Jens Petersen, and Klaus H Maier-Hein. nnu-net: a self-configuring method for deep learning-based biomedical image segmentation. *Nature Methods*, 18(2):203–211, 2021.
- Chen-Yu Lee, Saining Xie, Patrick Gallagher, Zhengyou Zhang, and Zhuowen Tu. Deeply-supervised nets. In *Artificial Intelligence and Statistics (AISTATS)*, pages 562–570, 2015.
- Tsung-Yi Lin, Piotr Dollár, Ross Girshick, Kaiming He, Bharath Hariharan, and Serge Belongie. Feature pyramid networks for object detection. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 2117–2125, 2017.

- Bjoern H Menze, Andras Jakab, Stefan Bauer, Jayashree Kalpathy-Cramer, Keyvan Farahani, Justin Kirby, Yvan Burren, Nicole Porz, Johannes Slotboom, Roland Wiest, et al. The multimodal brain tumor image segmentation benchmark (brats). *IEEE Transactions on Medical Imaging*, 34(10):1993–2024, 2014.
- Fausto Milletari, Nassir Navab, and Seyed-Ahmad Ahmadi. V-net: Fully convolutional neural networks for volumetric medical image segmentation. In *Fourth International Conference on 3D Vision (3DV)*, pages 565–571, 2016.
- David MW Powers. Evaluation: from precision, recall and f-measure to roc, informedness, markedness and correlation. *Journal of Machine Learning Technologies*, 2(1):37–63, 2011.
- Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. In *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, pages 234–241, 2015.
- John G Sled, Alex P Zijdenbos, and Alan C Evans. A nonparametric method for automatic correction of intensity nonuniformity in mri data. *IEEE Transactions on Medical Imaging*, 17(1):87–97, 1998.
- Abdel Aziz Taha and Allan Hanbury. Metrics for evaluating 3d medical image segmentation: analysis, selection, and tool. *BMC Medical Imaging*, 15(1):29, 2015.
- Naftali Tishby, Fernando C Pereira, and William Bialek. The information bottleneck method. *arXiv preprint physics/0004057*, 2000.
- Cédric Villani. *Optimal Transport: Old and New*, volume 338. Springer Science & Business Media, 2009.
- Tongzhou Wang and Phillip Isola. Understanding contrastive representation learning through alignment and uniformity on the hypersphere. In *International Conference on Machine Learning (ICML)*, pages 9929–9939, 2020.
- Haiyu Wu, Randall Balestriero, and Morgan Levine. Visreg: Variance-invariance-sketching regularization for japa training. *arXiv preprint arXiv:2606.02572*, 2026.
- Yuxin Wu and Kaiming He. Group normalization. In *European Conference on Computer Vision (ECCV)*, pages 3–19, 2018.

A Comprehensive Mathematical Notation and Extended Proofs

A.1 Structured Notation Reference Tables

Table 4: Comprehensive Mathematical Notation Reference.

Symbol	Mathematical Domain	Definition / Description
$\mathbf{X}$	$\mathbb{R}^{B \times 4 \times 128 \times 128 \times 128}$	Normalized 4-channel 3D input tensor
$\mathbf{Y}$	$\{0, 1\}^{B \times 1 \times 128 \times 128 \times 128}$	Ground-truth Whole Tumor (WT) binary mask
$\mathbf{E}_{\text{patches}}^{3D}$	$\mathbb{R}^{B \times 512 \times 384}$	3D Patch embedding tokens ($8^3 = 512$ tokens, $D = 384$)
$\mathbf{E}_{\text{pos}}^{3D}$	$\mathbb{R}^{1 \times 512 \times 384}$	Separable 3D sinusoidal coordinate positional embeddings
N_{ctx}	$\mathbb{N} = 192$	Constant visible context patch count (37.5% of volume)
N_{tgt}	$\mathbb{N} \approx 108$	Total target patch count across 4 cuboids (3^3 each)
E_{θ}	Neural Network	Online 3D ViT Context Encoder
$E_{\bar{\theta}}$	Neural Network	EMA Target Teacher Encoder (3D I-JEPA)
P_{ϕ}	Neural Network	3D JEPA Predictor (4 layers, $D_{\text{pred}} = 192$)
$\mathcal{T}_{\text{EP}}$	$\mathbb{R}_{\geq 0}$	Epps–Pulley Gaussianity test statistic with scale factor N
W_2^2	$\mathbb{R}_{\geq 0}$	1D Sliced-Wasserstein Distance against Gaussian quantiles
$\text{erank}(Z)$	$[1, D]$	Effective Rank using covariance eigenvalues S_k^2
$\bar{S}_{\text{centered}}$	$[-1, 1]$	Centered pairwise cosine similarity
d_{HD95}^{3D}	$\mathbb{R}_{\geq 0}$ (mm)	Exact 3D 95th Percentile Hausdorff Distance

A.2 Extended Theoretical Proofs

Lemma 1 (Cramér–Wold Hyperspherical Gaussianity Identification). *Let Z be a random vector in $\mathbb{R}^D$. If for every unit vector $u \in \mathbb{S}^{D-1}$, the linear projection $u^\top Z$ is distributed as a univariate standard normal $\mathcal{N}(0, 1)$, then Z is a standard multivariate Gaussian: $Z \sim \mathcal{N}(0, I_D)$.*

Proof. The characteristic function of Z is $\phi_Z(t) = \mathbb{E}[\exp(it^\top Z)]$. For any $t \in \mathbb{R}^D \setminus \{0\}$, let $r = \|t\|_2 > 0$ and $u = t/r \in \mathbb{S}^{D-1}$. Then:

$$\phi_Z(t) = \mathbb{E}[\exp(ir(u^\top Z))] = \phi_{u^\top Z}(r) \quad (35)$$

Since $u^\top Z \sim \mathcal{N}(0, 1)$, its characteristic function is $\phi_{u^\top Z}(r) = \exp(-r^2/2) = \exp(-\|t\|_2^2/2)$. This coincides uniquely with the characteristic function of $\mathcal{N}(0, I_D)$, establishing $Z \sim \mathcal{N}(0, I_D)$ by uniqueness of the Fourier transform. $\square$

Proposition 5 (Closed-Form 1D 2-Wasserstein Distance via Sorting). *Let $\mu_N = \frac{1}{N} \sum_{i=1}^N \delta_{p_i}$ be the empirical measure of 1D projections and let $\nu = \mathcal{N}(0, 1)$ with cumulative distribution function $\Phi(x)$. The 2-Wasserstein distance between μ_N and ν is computed in closed form in $\mathcal{O}(N \log N)$ time:*

$$W_2^2(\mu_N, \nu) = \frac{1}{N} \sum_{i=1}^N (p_{(i)} - q_i^*)^2 \quad (36)$$

where $p_{(1)} \leq \dots \leq p_{(N)}$ are the sorted projections and $q_i^* = \Phi^{-1}\left(\frac{i-0.5}{N}\right) = \sqrt{2} \operatorname{erf}^{-1}\left(2\frac{i-0.5}{N} - 1\right)$.

Proof. By classical optimal transport theory in one dimension [Villani, 2009, Bonneel et al., 2015], the optimal transport map between two 1D probability measures with cumulative distribution functions F and G is the unique monotone rearrangement $T(x) = G^{-1}(F(x))$. The continuous 2-Wasserstein distance is:

$$W_2^2(\mu, \nu) = \int_0^1 |F^{-1}(t) - G^{-1}(t)|^2 dt \quad (37)$$

For an empirical measure μ_N with order statistics $p_{(1)} \leq \dots \leq p_{(N)}$, the empirical inverse CDF is piecewise constant: $F^{-1}(t) = p_{(i)}$ for $t \in [\frac{i-1}{N}, \frac{i}{N})$. Midpoint discretization of $G^{-1}(t) = \Phi^{-1}(t)$ yields $q_i^* = \Phi^{-1}\left(\frac{i-0.5}{N}\right)$. Integrating across the N equal intervals yields the exact discrete sum $\frac{1}{N} \sum_{i=1}^N (p_{(i)} - q_i^*)^2$. Sorting requires $\mathcal{O}(N \log N)$ operations. $\square$

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [\[Yes\]](#)

Justification: The main claims are formulated explicitly in Section 1 (Core Research Questions RQ1–RQ4, Summary of Contributions) and supported by empirical benchmarks in Section 4 and formal proofs in Section 2 and Appendix A.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: Section 2, Section 3, and Appendix A disclose all architectural parameters, 3D BFS masking parameters, optimization schedules, loss weights, and deterministic seeding protocols needed for exact reproduction. The complete codebase is packaged in `thesis_3d`.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not include experiments.
- If the paper includes experiments, a [\[No\]](#) answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) The code is made available as part of the submission or is promised to be released upon acceptance.
 - (b) The data is made available as part of the submission or is promised to be released upon acceptance.
 - (c) A pipeline for recreating the data is made available as part of the submission or is promised to be released upon acceptance.

3. Open access to data and code

Question: Does the paper provide open access to data and code?

Answer: [\[Yes\]](#)

Justification: The code repository `thesis_3d` is fully standalone with automated pytest suites, reproducible configs, CLI runners, and Kaggle runner notebooks. The BraTS 2024 GLI challenge dataset is publicly available through Synapse.

4. Experimental setting/details

Question: Does the paper specify all the training and evaluation details?

Answer: [\[Yes\]](#)

Justification: Section 3 details optimizer choices (AdamW), learning rates, warmup schedules, weight decays, batch sizes, mixed-precision settings, and evaluation metrics.

5. Experiment statistical significance

Question: Does the paper report error bars or the statistical significance of the results?

Answer: [Yes]

Justification: All volumetric metrics (3D Dice, IoU, cKDTree HD95) are reported with sample means and standard deviations across the held-out test cohort.

6. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources needed to reproduce the experiments?

Answer: [Yes]

Justification: Section 3.2 and Section 4.1 report hardware specifications (NVIDIA Tesla T4 16GB / Apple Silicon), training durations per epoch, peak VRAM allocations, and forward inference latencies (in ms per 128^3 volume).

7. Code of ethics

Question: Does the research conform with the NeurIPS Code of Ethics?

Answer: [Yes]

Justification: The research uses openly published, fully de-identified clinical MRI challenge data (BraTS 2024 GLI) and strictly conforms with institutional ethics and the NeurIPS Code of Ethics.

8. Broader impacts

Question: Does the paper discuss potential positive and negative societal impacts?

Answer: [Yes]

Justification: Section 5.3 outlines the clinical utility of label-efficient 3D segmentation in resource-constrained radiotherapy planning, while addressing clinical safety boundaries against unsupervised off-label diagnostic deployment.

9. Safeguards

Question: Does the paper describe safeguards for responsible release of models?

Answer: [N/A]

Justification: The models are discriminative volumetric encoders dedicated to brain glioma MRI segmentation and pose no dual-use or disinformation risks.

10. Licenses for existing assets

Question: Are creators of existing assets properly credited and licenses respected?

Answer: [Yes]

Justification: Open benchmarks (BraTS 2024 GLI) and open-source libraries (PyTorch, MONAI, NiBabel) are cited with proper attributions.

11. New assets

Question: Are new assets introduced in the paper well documented?

Answer: [N/A]

Justification: The paper utilizes existing public benchmark datasets.

12. Crowdsourcing and research with human subjects

Question: Does the paper involve crowdsourcing or research with human subjects?

Answer: [N/A]

Justification: No crowdsourcing or human subject experiments were performed.

13. Institutional review board (IRB) approvals

Question: Does the paper describe IRB approvals for research with human subjects?

Answer: [N/A]

Justification: The research exclusively utilizes previously published, de-identified public challenge datasets with pre-existing ethics approvals.

14. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs?

Answer: [N/A]

Justification: Large language models were not used as a methodological component of the algorithmic design or experimental pipeline.